

# COMP34711 Week 3

## Part 2 Querying and ranking

Goran Nenadic

with examples from the IIR book

# Outline

- The Boolean model
- Ranked retrieval
  - Vector space model for IR
  - $tf*idf$
  - Ranking documents
- Measuring the quality of IR
  - Standard metrics
- Enhancing IR
  - User behaviour

# Boolean queries


- The simplest query model: find all documents from the collection that fully match the query
  - Binary outcome for each document: **yes/no**
- Use operators
  - AND (set intersection)
  - OR (set union)
  - NOT (set difference, complement)
- E.g.
  - drug AND approach** = all documents that contain both **query words**
  - drug OR approach** = all documents that either of the **query words**
  - drug AND NOT approach** = all documents that contain **drug** but do not contain **approach**

# Boolean query processing: **AND**

- Consider processing the query:

## *Brutus AND Caesar*


1. Locate *Brutus* in the Dictionary
  - Retrieve its postings
2. Locate *Caesar* in the Dictionary
  - Retrieve its postings



3. “AND-merge” the two postings:

# Boolean query processing: AND

- “AND-merge”:
  - walk through the two postings simultaneously (moving/progressing through the postings with lower docID) until there is no possible matches
  - crucial: postings are sorted by docID.



# Boolean model: **pros**

- Simple model: everything is considered as a set of terms
- Easy/efficient to implement
- Precise
  - Document either matches or does not match
- Widely used for commercial, legal retrieval and for specialist searches
  - Long, precise queries
- Works well when we know what we want: user feels in control

# Boolean model: **cons**

- Users are not good in formulating queries: possible confusion with natural language
  - cats OR dogs
  - cats OR dogs AND NOT horses
- “feast or famine”
  - AND gives too few or no results  
(the more ANDs, the smaller the result set)
  - OR gives too many
- Basic Boolean expressions too limiting for information needs?

# Boolean model: **cons**

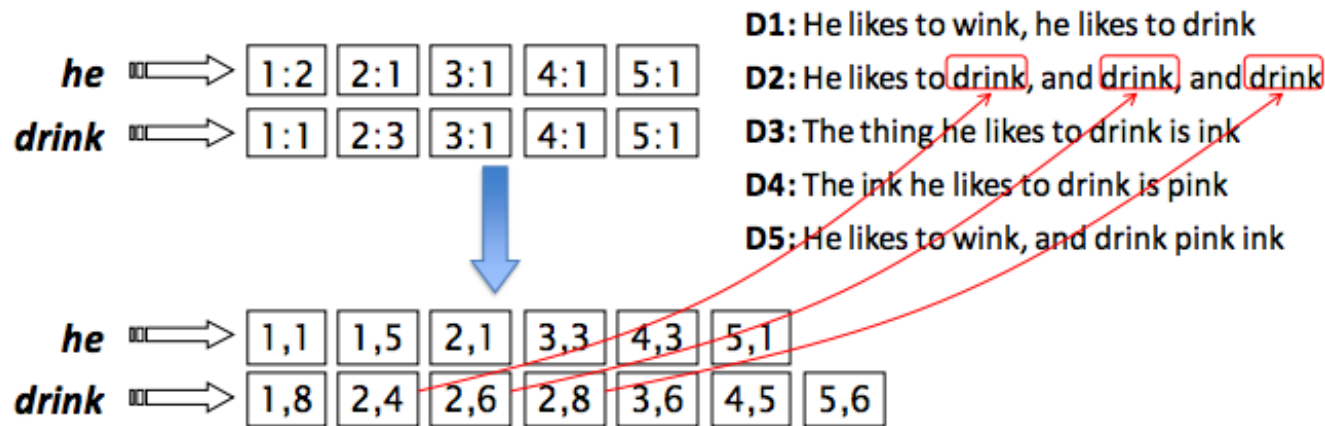
- No relevance ordering/ranking of results
  - In principle, all retrieved documents are good
    - But note: the BoW model does not use word order in documents
  - How to ‘read’ them if there are too many?
  - Some notion of additional relevance could be added:
    - date reverse order of document creation?
    - the frequency of query terms in matched documents?
    - proximity of query terms in documents?



# Extended Boolean model

- Proximity operators**

- Embed term **positions** to the inverted index (proximity index)



- Queries can then refer to these, e.g.

- */n* (e.g. */3* – within 3 words)
- */s* = in same sentence, */p* in same paragraph
- *+s* = term1 must precede term2 in same sentence

# Ranked retrieval


- In Boolean query model
  - Documents match or don't match
  - Results are unranked
- Documents that are 'close' are not retrieved
  - social AND worker AND union**
    - Will not retrieve a document that has **social** and **worker** in it if **union** is not mentioned, although it may mention **UCU** or **TUC**
    - Boolean querying is too rigid

# Ranked retrieval

- Introducing similarity and ranking of matched documents
  - Attempt more than exact matching queries with docs
  - Score each document to say how well it matches query
    - E.g. assign a real number score in range 0..1
  - Typically aim to get top  $K$  (10?) ones correct
    - As user will likely not look much further
- Idea: use vector representation for both documents and queries, and calculate their similarity

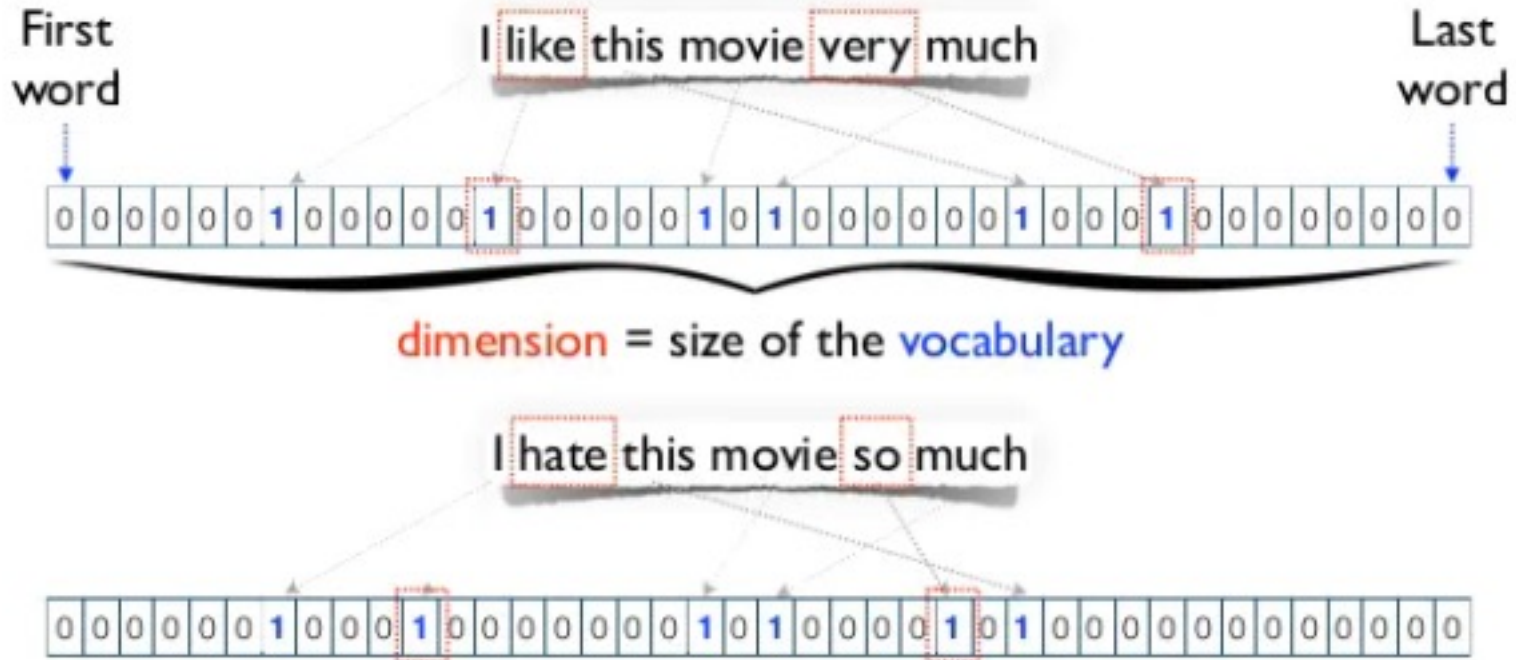
# Vector representations

- Represent both documents and queries as vectors in the same space
- Then, rank (all) documents according to their proximity to the given query in that space
  - Rank more relevant documents higher than less relevant documents
  - Recall: We do this because we want to get away from the “you’re-either-in-or-out” Boolean model.



# Weights

- What weight to use to represent terms that appear in documents/queries?



# Term frequency (tf)

- Term frequency  $tf_{t,d}$  of term  $t$  in document  $d$  is defined as the number of times that  $t$  occurs in  $d$ .
- How to use tf when computing query-document match scores?
- Raw term frequency is not what we want:
  - Document with 10 occurrences of the term is more relevant than a document with 1 occurrence of the term
  - *But not 10 times more relevant*
- Relevance does not increase proportionally with term frequency

NB: frequency = count in IR

# Document frequency (df)

- Document frequency = in how many documents term appears
- Consider term in query that is rare in collection  
(e.g., *arachnocentric*)
- A document containing this term is very likely to be relevant to a query that contains *arachnocentric*
  - We want a high weight for rare terms  
(like *arachnocentric*)
- Rare terms are more informative and discriminative than frequent terms for IR

# Document frequency (df)

- Consider a query term that is frequent in collection (e.g., *high*, *increase*, *line*)
- A document containing such a term is more likely to be relevant than a document that doesn't
- But it's not a sure indicator of relevance
  - For frequent terms, we want high positive weights, but lower weights than for rare terms



# Collection vs. document frequency

- The collection frequency of *t* is the number of occurrences of *t* in the collection, counting multiple occurrences (within the same document)
- Example:

Word	Collection frequency	Document frequency
<i>insurance</i>	10440	3997
<i>try</i>	10422	8760

- Which word is a “better” search term (and should get a higher weight)?

# Inverse document frequency

- Document frequency ( $df_t$ ) is the number of documents that contain term  $t$ 
  - $df_t$  is an inverse measure of the “informativeness” of  $t$
  - $df_t \leq N$
- We define the idf (inverse document frequency) of  $t$  by

$$idf_t = \log_{10} (N/df_t)$$

- We use  $\log(N/df_t)$  instead of  $N/df_t$  to “dampen” the effect of idf
- The base of the log is immaterial.

# idf example

Collection of  $N = 1,000,000$  documents

term	$df_t$	$idf_t$
Calpurnia	1	6 (= $\log(1M/1)$ )
animal	100	4 (= $\log(1M/100)$ )
sunday	1,000	3 (= $\log(1000)$ )
fly	10,000	2 (= $\log(100)$ )
under	100,000	1 (= $\log(10)$ )
the	1,000,000	0 (= $\log(1)$ )

- There is one idf value for each term  $t$  in a collection

$$idf_t = \log_{10} (N/df_t)$$

$$idf_{fly} = \log_{10} (1,000,000/10,000) = \log_{10} 100 = 2$$

# tf.idf weighting

- The **tf-idf** weight of a term is the product of its **tf** weight and its **idf** weight

$$\text{tf.idf}_{t,d} = (1 + \log_{10} \text{tf}_{t,d}) \times \log_{10} (N / \text{df}_t)$$

- Best known weighting scheme in information retrieval
  - Note alternative notations: tf-idf, tf x idf, tf\*idf, tfidf
- **Increases**
  - with the number of occurrences within a document
  - with the rarity of the term in the collection

frequency

# Term-document count matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	157	73	0	0	0	0
Brutus	4	157	0	1	0	0
Caesar	232	227	0	2	1	1
Calpurnia	0	10	0	0	0	0
Cleopatra	57	0	0	0	0	0
mercy	2	0	3	5	5	1
worser	2	0	1	1	1	0

Each document is a count vector in  $\mathbb{N}^v$


Binary  $\rightarrow$  count  $\rightarrow$  tf\*idf weight matrix

	Antony and Cleopatra	Julius Caesar	The Tempest	Hamlet	Othello	Macbeth
Antony	5.25	3.18	0	0	0	0.35
Brutus	1.21	6.1	0	1	0	0
Caesar	8.59	2.54	0	1.51	0.25	0
Calpurnia	0	1.54	0	0	0	0
Cleopatra	2.85	0	0	0	0	0
mercy	1.51	0	1.9	0.12	5.25	0.88
worser	1.37	0	0.11	4.15	0.25	1.95

Each document is now represented by a real-valued vector of tf\*idf weights  $\in \mathbb{R}^{|V|}$


# Vector Space Model for IR

- Documents and Queries are presented as vectors
- $\text{Match}(Q,D)$  = distance between vectors
- Which distance to use?
  - Euclidean Distance?
    - Distance between the endpoints of the two vectors
    - Large for vectors of diff. lengths
  - Angle between the document and the query
    - Use cosine of the angle



# Cosine distance

- Need to normalise for length to ensure fair comparison
  - Long and short documents then have comparable weights
- Dividing a vector by its **norm** makes it a unit (length) vector (on surface of unit hypersphere)





# Cosine distance

- For ***length-normalized vectors***, cosine similarity is the dot (or scalar) product:

$$\cos(\vec{q}, \vec{d}) = \vec{q} \cdot \vec{d} = \sum_{i=1}^{|V|} q_i d_i$$

- Easy and efficient to calculate
  - Note that we need only the values for terms that appear in both the document and query.
- Use cosine values to rank the document based on their similarity (i.e. distance) to the query

# Document ranking for a query

$$\text{Score}(q, d) = \sum_{t \in q \cap d} \text{tf.idf}_{t,d}$$

- Score of a document  $d$  is sum over all tf.idf weights of each query term  $t$  found in  $d$
- Indexing: for each term and document calculate  $\text{tf} \cdot \text{idf}_{t,d}$
- Typical output: a list of documents ranked according to score (q,d):

1	0	710	0	0.9234	0
1	0	213	0	0.7678	0
1	0	103	0	0.6761	0
1	0	13	0	0.6556	0
1	0	501	0	0.4301	0

Return top K documents

Query id

document id

score

# Summary of the steps

- Pre-process each document
  - tokenisation, stop-words removal (plus maybe some normalization – stemming, spelling corrections)
  - decide what will be index terms and calculate their tf.idf
- Represent each document as a weighted tf.idf vector
- Represent the query as a weighted tf.idf vector
- Compute the (cosine) similarity score for the query vector and each document vector
- Rank documents with respect to the query by score
- Return the top K to the user