

🔑 2bb083470d ▾

⋮

ProductReviewClassifier

/

Archive

/

Task2.ipynb

🐙

Boyko Borisov

Migrate repository to personal git

🕒

👤 0 contributors

779 lines (779 sloc) | 34.9 KB

⋮

```
In [1]: import nltk
import string
import os
import copy
import numpy as np
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import re
from random import sample, shuffle
from sklearn.cluster import KMeans, AgglomerativeClustering
device = torch.device('cuda' if torch.cuda.is_available() else 'cpu')
```

```
In [2]: nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

```
Out[2]: True
```

```
In [3]: # I used google collab for training my model, uncomment the lines below to
# connect to google drive in google colab
# NB corpus_root SHOULD BE CHANGED TO MATCH THE CORPUS PATH ON THE SPECIFIC MACHINE

# from google.colab import drive
# drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount()

```
In [4]: stop_words = set(nltk.corpus.stopwords.words("english"))
stemmer = nltk.SnowballStemmer("english", ignore_stopwords = False)
# Folder path where corpus root should be
corpus_root = r"/content/drive/MyDrive/cw2/product_reviews"
file_pattern = r"*.txt"
original_corpus = nltk.corpus.PlaintextCorpusReader(corpus_root, file_pattern)
print(original_corpus.fileids())
```

```
['Canon_PowerShot_SD500.txt', 'Canon_S100.txt', 'Diaper_Champ.txt', 'Hitachi_route
'norton.txt']
```

```
In [5]: # Core utility function for document cleaning
# Works recursively, split the text into sentences/review, then for each sentence
def process_doc(text, remove_punctuation, case_fold, stem,
               remove_stopwords, remove_short_tokens, tokenize_by, stem_blacklist,
               remove_nonalphabetical = False):

    if (tokenize_by == "sentence"):
        sentences = nltk.RegexpTokenizer("##", gaps = True).tokenize(text)
        sentences = [process_doc(sentence, remove_punctuation, case_fold, stem,
                                remove_stopwords, remove_short_tokens, "words", stem_blacklist,
                                remove_nonalphabetical) for sentence in sentences]
        return sentences
    if (tokenize_by == "reviews"):
        reviews = nltk.RegexpTokenizer("W[ t W]", gaps = True).tokenize(text)
        reviews = [process_doc(review, remove_punctuation, case_fold, stem,
                                remove_stopwords, remove_short_tokens, "words", stem_blacklist,
                                remove_nonalphabetical) for review in reviews]
        return reviews
```

```

if (tokenize_by == "words"):
    words = nltk.TreebankWordTokenizer().tokenize(text)
    if (remove_punctuation):
        words = [w for w in words if w not in string.punctuation and w != "..."]
        words = [w.strip("'") for w in words]
        words = [w.strip(".") for w in words]
    if (case_fold):
        words = [w.lower() for w in words]
    if (remove_short_tokens):
        words = [w for w in words if len(w) > 2]
    if (stem):
        words = [w if w in stem_blacklist else stemmer.stem(w) for w in words]
    if (remove_stopwords):
        words = [w for w in words if w not in stop_words and w != "n't"]
    if (remove_punctuation):
        words = [w for w in words if w not in string.punctuation and w != "..."]
    if (remove_nonalphabetical):
        words = [w for w in words if w.isalpha()]
    return words

def process_corpus(corpus, remove_punctuation:bool, case_fold:bool, stem:bool,
                  remove_stopwords:bool, remove_short_tokens, tokenize_by:str, re
docs = [word for fileid in corpus.fileids()
        for word in process_doc(corpus.raw(fileid), remove_punctuation, case
                              stem, remove_stopwords, remove_short_tokens,
                              tokenize_by, remove_nonalphabetical)]

return docs

def most_frequent(words, n, should_print):
    freqDist = nltk.FreqDist(words)
    most_common = freqDist.most_common(n)
    if (should_print):
        i = 1
        for (w, count) in most_common:
            print(i, w, count)
            i += 1
    return most_common

```

In [6]:

```

def get_all_sentences_cleaned(corpus_filepath):
    corpus = nltk.corpus.PlaintextCorpusReader(corpus_filepath, file_pattern)
    out = []
    for fileid in corpus.fileids():
        sentences = process_doc(corpus.raw(fileid), True, True, True, True, True,
                              out.extend(sentences)
    return out

# partitions corpus into sentiments, and cleans the text
def get_all_sentiments_cleaned(corpus_filepath, stemming, stop_words, remove_mix
corpus = nltk.corpus.PlaintextCorpusReader(corpus_filepath, file_pattern)
# pattern used to match sentiments
pattern = re.compile(r"(([a-z -]*W[W-W+][0-9]W),? ?)+#[^(W[)]+)"
sentiments = []
for file_id in corpus.fileids():
    text = corpus.raw(file_id)
    text = re.sub("W[[a-z]+W]", "", text)
    text = pattern.findall(text)
    for sentiment in text:
        sentiment_parsed = sentiment[0]
        # Find all labels for whether a sentiment is positive or negative
        matches = re.findall("W[W+W-][0-9]W", sentiment_parsed)
        score = 0
        has_positive = False

```

```

        has_negative = False
        for match in matches:
            score += int(match[1:-1])
            if match[1] == "+":
                has_positive = True
            if match[1] == "-":
                has_negative = True
        # if the sum of all scores is 0 discard the sample, since we are doing binary
        # classification, optionally remove all sentiments with mixed labels
        if remove_mixed_sentiments and has_positive and has_negative:
            continue
        if (score == 0): continue
        if (score < 0): score = 0
        if (score > 0): score = 1
        sentiment_parsed = process_doc(sentiment_parsed, True, True, stemming, stop_words)
        if (len(sentiment_parsed[: -1]) < 2): continue
        sentiments.append((sentiment_parsed[: -1], score))
    return sentiments

def generate_word_to_idx_and_idx_to_word(corpus):
    word_to_idx = {}
    idx_to_word = {}
    i = 0
    for sentence in corpus:
        for word in sentence[0]:
            if (word not in word_to_idx):
                word_to_idx[word] = i
                idx_to_word[i] = word
                i += 1
    return (word_to_idx, idx_to_word)

def get_context_window_tuples(word_to_idx, sentences, window, key_words):
    tuples = []
    for sentence in sentences:
        for i in range(window, len(sentence) - window):
            # if sentence[i] in key_words:
            context = []
            middle_word = word_to_idx[sentence[i]]
            for j in range(i - window, i + window + 1):
                if i != j:
                    context.append(word_to_idx[sentence[j]])
            tuples.append((context, word_to_idx[sentence[i]]))

    return tuples

def get_skipgrams(sentiments, window):
    word = []
    context = []
    for sentiment in sentiments:
        sentence = sentiment[0]
        for i in range(len(sentence)):
            cont = [sentence[idx] for idx in range(max(0, i - window), min(len(sentence), i + window + 1))]
            word.append(sentence[i])
            context.extend(cont)
    return (word, context)

def get_batches(words, contexts, batch_size):
    shuffled_idx = sample(range(0, len(words)), len(words))
    batches = []

    batch_word, batch_context = [], []
    for i in range(len(words)):
        word = words[shuffled_idx[i]]
        context = contexts[shuffled_idx[i]]
        batch_word.append(word)
        batch_context.append(context)
        if len(batch_word) == batch_size:
            batches.append((batch_word, batch_context))
            batch_word, batch_context = [], []
    if len(batch_word) > 0:
        batches.append((batch_word, batch_context))
    return batches

```

```

        idx = shuffled_idxes[i]
        batch_word.append(words[idx])
        batch_context.append(contexts[idx])
        if (i + 1) % batch_size == 0 or i + 1 == len(words):
            batches.append((
                torch.from_numpy(np.array(batch_word)),
                torch.from_numpy(np.array(batch_context))
            ))
            batch_word, batch_context = [], []
    return batches

def get_x_tensors(x_y_tuples):
    tensors = []
    for tuple in x_y_tuples:
        tensors.append(torch.tensor(tuple[0], dtype=torch.long))
    return tensors

def get_y_tensors(tuples, num_classes):
    tensors = []

    for tuple in tuples:
        tensors.append(F.one_hot(torch.tensor(tuple[1]), num_classes=num_classes))
    return tensors

def get_sentiments_as_word_idxes(sentiments, word_to_idx):
    return [[word_to_idx[word] for word in words], label] for (words, label) in

```

In [7]:

```

# Function to split data into K folds
def k_fold_partitioning(sentiments, k, should_shuffle):
    # shuffle
    # we do not shuffle when we need to compare the results of experiments
    if (should_shuffle):
        shuffle(sentiments)
    folds = []
    # determine fold size
    partition_step = len(sentiments) // k
    remainders = len(sentiments) % k
    start = 0
    # append to each fold
    for i in range(k):
        if (remainders > 0):
            folds.append(sentiments[start : start + partition_step + 1])
            start += partition_step + 1
            remainders -= 1
        else:
            folds.append(sentiments[start : start + partition_step])
            start += partition_step
    return folds

# partition the data into two - the i-th fold and the rest
def split_training_testing_from_k_folds(i, folds):
    # To ensure that not tampering is done
    testing = copy.deepcopy(folds[i])

    training = []
    for j in range(i):
        training.extend(folds[j])
    for j in range(i + 1, len(folds)):
        training.extend(folds[j])
    return (training, testing)

def to_tensors(sentiments):
    # Convert the sentiments to tensors

```

```

shuffle(sentiments)
start = 0
batches = [(torch.from_numpy(np.array(sentiment[0])), torch.from_numpy(np.array(
return batches

# y_hat is a tensor output of a sigmoid (y_hat between: [0, 1])
def get_binary_accuracy(y_hat, y, verbose=False):
    # if y_hat <= 0.5: rounded = 0 else: rounded = 1
    rounded = torch.round(y_hat)
    correct = (rounded == y).float()
    if verbose:
        print("y_hat:", y_hat.data)
        print("y:", y.data)
        print("rounded:", rounded.data)
        print("correct: ", correct.data)
    return correct

```

In [8]:

```

class CNN(nn.Module):
    def __init__(self, vocab_size, n_filters, embedding_dim = None, padding_idx = None):
        super().__init__()

        if (embedding_weights != None):
            self.embedding = nn.Embedding.from_pretrained(embedding_weights, freeze=True)
            embedding_dim = embedding_weights.size()[1]
        else:
            self.embedding = nn.Embedding(vocab_size, embedding_dim, padding_idx = padding_idx)

        self.conv_0 = nn.Conv2d(in_channels = 1,
                                out_channels = n_filters,
                                kernel_size = (3, embedding_dim))
        self.conv_1 = nn.Conv2d(in_channels = 1,
                                out_channels = n_filters,
                                kernel_size = (4, embedding_dim))
        self.conv_2 = nn.Conv2d(in_channels = 1,
                                out_channels = n_filters,
                                kernel_size = (5, embedding_dim))

        self.fc = nn.Linear(3 * n_filters, 1)
        self.dropout = nn.Dropout(dropout_rate)
        self.sigmoid = nn.Sigmoid()

    def forward(self, text, training = False):

        embedding = self.embedding(text)
        #embedding = [len(text) x embedding_size]

        embedding = embedding.unsqueeze(1)
        embedding = embedding.unsqueeze(1)
        embedding = embedding.permute(1, 2, 0, 3)

        convded_0 = F.relu(self.conv_0(embedding).squeeze(3))
        convded_1 = F.relu(self.conv_1(embedding).squeeze(3))
        convded_2 = F.relu(self.conv_2(embedding).squeeze(3))
        #convded_n = [len(text) - kernel_size x number of filters]

        pooled_0 = F.max_pool1d(convded_0, convded_0.shape[2]).squeeze(2)
        pooled_1 = F.max_pool1d(convded_1, convded_1.shape[2]).squeeze(2)
        pooled_2 = F.max_pool1d(convded_2, convded_2.shape[2]).squeeze(2)

        concat = torch.cat((
            pooled_0,
            pooled_1,
            pooled_2,

```

```

        pooled_2)
        , dim = 1)
    # Apply dropout only when training
    if (training):
        concat = self.dropout(concat)

    #concat = [len(text) - kernel_size x number of filters]
    return self.sigmoid(self.fc(concat))

```

In [9]:

```

def train_and_eval(num_filters, embedding_size, epochs, batch_size, learning_rate):
    K = len(folds)
    # accuracies will contain the accuracies for each fold for each epoch
    accuracies = np.zeros((K, epochs))
    for k in range(K):
        if (verbose):
            print("FOLD: ", k + 1)
        (training_data, testing_data) = split_training_testing_from_k_folds(k, folds)
        model = CNN(vocab_size + 1, num_filters, embedding_size, padding_idx, dropout)
        model = model.to(device)
        optimizer = optim.Adam(model.parameters(), learning_rate)
        loss_fn = nn.BCELoss()
        loss_fn = loss_fn.to(device)
        for epoch in range(epochs):
            if (verbose):
                print("EPOCH:", epoch + 1)
            acc = 0
            total_loss = 0
            n = 0
            for sample in to_tensors(training_data):
                optimizer.zero_grad()
                (sentiment, label) = sample
                sentiment = sentiment.to(device)
                label = label.to(device)
                y_hat = model(sentiment, training = True).squeeze()
                acc += get_binary_accuracy(y_hat, label, 1 == 0)
                loss = loss_fn(y_hat, label.float())
                total_loss += loss
                loss.backward()
                optimizer.step()
                n += 1
            if (verbose):
                print("TRAINING: accuracy", (acc / n).item(), "total loss", total_loss.item())
            with torch.no_grad():
                accuracy = 0
                for (sentiment, label) in to_tensors(testing_data):
                    sentiment = sentiment.to(device)
                    label = label.to(device)
                    accuracy += get_binary_accuracy(model(sentiment), label)
                accuracies[k][epoch] = (accuracy / len(testing_data)).item()
            if (verbose):
                print("EPOCH VALIDATION ACCURACY", accuracy.item())

    # epoch averages contains the average accuracy for each epoch accross all folds
    epoch_averages = np.mean(accuracies, axis=0)
    best_epoch = np.argmax(epoch_averages)
    return "Best epoch:", best_epoch + 1, "with an average", epoch_averages[best_epoch]

```

In [10]:

```

def run_experiment(num_filters, embedding_size, epochs, batch_size, learning_rate):
    # sentiments - a list of tuples, tuple[0] is the cleaned text of a sentiment, tuple[1] is the sentiment label
    sentiments = get_all_sentiments_cleaned(corpus_root, stemming, stopword_removal)
    (word_to_idx, idx_to_word) = generate_word_to_idx_and_idx_to_word(sentiments)
    # tuple[0] in sentiments becomes a list of ints, each int represents a token, w

```

```

sentiments = get_sentiments_as_word_idxes(sentiments, word_to_idx)
PADDING_STR = ""
PADDING_IDX = len(word_to_idx)
idx_to_word[PADDING_IDX] = PADDING_STR
word_to_idx[PADDING_STR] = PADDING_IDX
vocab_size = len(idx_to_word)
# The filter size of the CNN is 5, all shorter texts than that need padding
for sentiment in sentiments:
    while (len(sentiment[0]) < 5):
        sentiment[0].append(PADDING_IDX)
k_folds = k_fold_partitioning(sentiments, 5, should_shuffle)

# training and evaluation
return train_and_eval(num_filters, embedding_size, epochs, batch_size, learning_rate,

```

```

In [11]: print("Removing mixed sentiments", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=False, stopwords=stopwords)

print("Keeping mixed sentiments", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=False, stopwords=stopwords)

```

Removing mixed sentiments ('Best epoch:', 5, 'with an average', 0.7084891080856324)  
Keeping mixed sentiments ('Best epoch:', 5, 'with an average', 0.6722772240638732)

```

In [12]: print("With stemming", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=True, stopwords=stopwords)

print("With stop words removal", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=False, stopwords=stopwords)

print("Baseline", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=False, stopwords=stopwords)

```

With stemming ('Best epoch:', 13, 'with an average', 0.6985148429870606)  
With stop words removal ('Best epoch:', 9, 'with an average', 0.6851130485534668)  
Baseline ('Best epoch:', 6, 'with an average', 0.7138613820075989)

```

In [13]: print("Dropout 0", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=False, stopwords=stopwords)

print("Dropout 0.25", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=False, stopwords=stopwords)

print("Dropout 0.50", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=False, stopwords=stopwords)

print("Dropout 0.75", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=False, stopwords=stopwords)

print("Dropout 0.85", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=False, stopwords=stopwords)

```

Dropout 0 ('Best epoch:', 20, 'with an average', 0.6866336584091186)  
Dropout 0.25 ('Best epoch:', 16, 'with an average', 0.6985148549079895)  
Dropout 0.50 ('Best epoch:', 9, 'with an average', 0.7054455399513244)  
Dropout 0.75 ('Best epoch:', 6, 'with an average', 0.705940580368042)  
Dropout 0.85 ('Best epoch:', 6, 'with an average', 0.7113861203193664)

```

In [16]: print("Dropout 0.95", run_experiment(num_filters = 100, embedding_size = 300, epochs = 20,
learning_rate=0.001, stemming=False, stopwords=stopwords)

```



Dropout 0.95 ('Best epoch:', 18, 'with an average', 0.6539603888988494)

```
In [14]: # Running for only 15 epochs to speed up the experiment
print("Filter number 50:", run_experiment(num_filters = 50, embedding_size = 300,
                                          learning_rate=0.001, stemming=False, stopw

print("Filter number 100", run_experiment(num_filters = 100, embedding_size = 300,
                                          learning_rate=0.001, stemming=False, stopw

print("Filter number 200", run_experiment(num_filters = 200, embedding_size = 300,
                                          learning_rate=0.001, stemming=False, stopw

print("Filter number 300", run_experiment(num_filters = 300, embedding_size = 300,
                                          learning_rate=0.001, stemming=False, stopw

print("Filter number 400", run_experiment(num_filters = 400, embedding_size = 300,
                                          learning_rate=0.001, stemming=False, stopw

Filter number 50: ('Best epoch:', 9, 'with an average', 0.6836633563041687)
Filter number 100 ('Best epoch:', 6, 'with an average', 0.6985148429870606)
Filter number 200 ('Best epoch:', 8, 'with an average', 0.6693069219589234)
Filter number 300 ('Best epoch:', 9, 'with an average', 0.6623762249946594)
Filter number 400 ('Best epoch:', 9, 'with an average', 0.6757425785064697)
```

```
In [20]: print("Embedding size 50", run_experiment(num_filters = 100, embedding_size = 50,
                                                    learning_rate=0.0005, stemming=False, stopw

print("Embedding size 100", run_experiment(num_filters = 100, embedding_size = 100,
                                           learning_rate=0.0005, stemming=False, stopw

print("Embedding size 200", run_experiment(num_filters = 100, embedding_size = 200,
                                           learning_rate=0.0005, stemming=False, stopw

print("Embedding size 300", run_experiment(num_filters = 100, embedding_size = 300,
                                           learning_rate=0.0005, stemming=False, stopw

print("Embedding size 400", run_experiment(num_filters = 100, embedding_size = 400,
                                           learning_rate=0.0005, stemming=False, stopw

Embedding size 50 ('Best epoch:', 11, 'with an average', 0.6975247383117675)
Embedding size 100 ('Best epoch:', 15, 'with an average', 0.6891089081764221)
Embedding size 200 ('Best epoch:', 10, 'with an average', 0.7064356327056884)
Embedding size 300 ('Best epoch:', 13, 'with an average', 0.7143564343452453)
Embedding size 400 ('Best epoch:', 12, 'with an average', 0.7099009871482849)
```

```
In [19]: print("Accuracy with best parameters:", run_experiment(num_filters = 100, embeddi
                                                    learning_rate=0.0005, stemming=False, stopw

Accuracy with best parameters: ('Best epoch:', 28, 'with an average', 0.7410756111
```